



PySpark Distributed Computing

Leveraging the Functional Model

Parallel Programming
(deterministic)

Concurrency
(non-deterministic)

Functional

Distr.

Bandwidth
Node failure
Connectivity

MapReduce
Spark

CAP theorem

Erlang
Akka
Pykka

Immutable data
Ref. transparency
Declarative
Streams

Local

Side effects
Low-level abstractions

Pool.map



Resource contention
Deadlocks
Thrashing

STM
Pypy
Twisted

Haskell
Erlang
Clojure, Scala

What is Apache *Spark*

- Fast and general engine for large-scale data processing
- Multi-stage in-memory primitives
- Supports Iterative Algorithms
- High-Level Abstractions
- Extensible; integrated stack of libraries



Spark Example

```
from pyspark import SparkContext
sc = SparkContext("local", "Log Analyzer")

rdd1 = sc.textFile('testfile1.txt')
rdd2 = sc.textFile('testfile2.txt')


def splitLine(line):
    key, val = line.split('|')
    return (key, int(val))

a = rdd1.filter(lambda line: not "ERROR" in line).map(splitLine)
b = rdd2.map(splitLine).reduceByKey(lambda a,b: a+b)

res = a.join(b)
#print res.toDebugString()
res.take(10)


[(u'key3', (50, 30)), (u'key1', (20, 60))]
```

Operator Graph






vs.



- Arbitrary operator graph
- Lazy eval of lineage graph => optimization
- Off-heap use of large memory
- Native integration with python



RDD

- Resilient Distributed Datasets are primary abstraction in Spark
- fault-tolerant collection
 - parallelized collections
 - hadoop datasets
- can be cached for reuse
- extensions (SchemaRDD)

Transformations

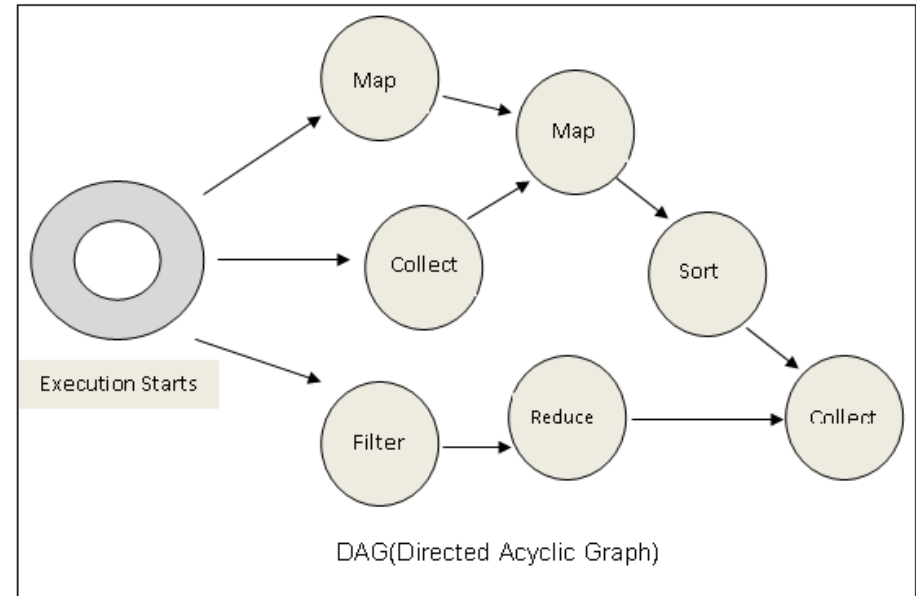
map()
filter()
flatMap()
mapPartitions()
sample()
union()
intersection()
distinct()
groupByKey()
reduceByKey()
aggregateByKey()
sortByKey()
join()
cogroup()
cartesian()
coalesce()
repartition()

Actions

reduce()
collect()
count()
take()
takeSample()
takeOrdered()
saveAsTextFile()
saveAsSequenceFile()
countByKey()
foreach()

Lifetime of an RDD

1. create from data
 - local collection
 - hadoop data set
2. lazily combine RDDs using transformations
 - `map()`
 - `join()`
 - etc.
3. call an RDD 'action' on it (`collect()`, `count()`, etc.) to "collapse" tree:
 1. Operator DAG is constructed
 2. Split into stages of tasks
 3. Launch tasks via cluster manager
 4. Execute tasks on local machines
4. store/consume results



Integrated Libraries

Spark
SQL

Spark
Streaming

MLlib
(machine
learning)

GraphX
(graph)

Apache Spark

Takeaways


Spark...

- feels like native python, very nice API
- adds awesome Distributed Computing and Parallel Programming capabilities to python
- comes with batteries included (SQL, GraphX, MLlib, Streaming, etc.)
- can be used from the start for exploratory programming



Getting Started

- Download Spark; ./bin/pyspark
- [docker-spark](#)
- [Spark on Amazon EMR](#)
- [Berkeley MOOC setup](#)
(vagrant, virtualbox, notebook)



Backups

Pure Functions

- $f: a \rightarrow b$
- Takes an “a” and returns a “b”
- Does not access global state and has no side-effects
- Function invocation can be substituted with the function body
- Can be used in an expression
- Can be “memoized”
- Is idempotent

Pure

- stateless
- no sequence, no time
- non-strict
- $x = 1+4$ (equality)
- “x” can be substituted by the expression
(referential transparency)
- idempotent
- expressions, algebra

Effects

- stateful
- fixed sequence, time
- strict
- $x := x + 1$ (assignment)
- “x” = changeable memory
“slot”

Pure functions by themselves are useless.


We want to interact with storage, network, screen etc.

We need both pure functions and (controlled, contained) effects

Immutable State

`append([1, 2, 3], 4) => [1, 2, 3, 4]`

- `[1, 2, 3]` remains unchanged
- Inherently thread-safe
- Can be shared freely
- “Everything is atomic”



Streams (Generators, Iterators)

Declarative

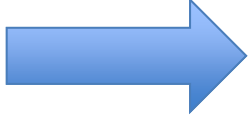
```
xs = [1, 2, 3];  
return xs.map(x => x+1);
```

Imperative

```
xs = [1, 2, 3];  
res = []  
for (int i = 0; i < 3; i++) {  
    res.append(xs[i] + 1);  
}  
return res;
```

Which do you think is easier to parallelize?

Stream Fusion

`xs`
`.map(x => x+1)`
`.map(y => y*2)`  `xs`
`.map(x => (x+1)*2)`

If functions are pure, we can

- combine
- reorder
- optimize the entire chain

If application is lazy, we can optimize across functions as well