

# Spark and Kafka for Data Analysis and Machine Learning

Cork JUG, 2019-02-19

Johannes Ahlmann




# About Me

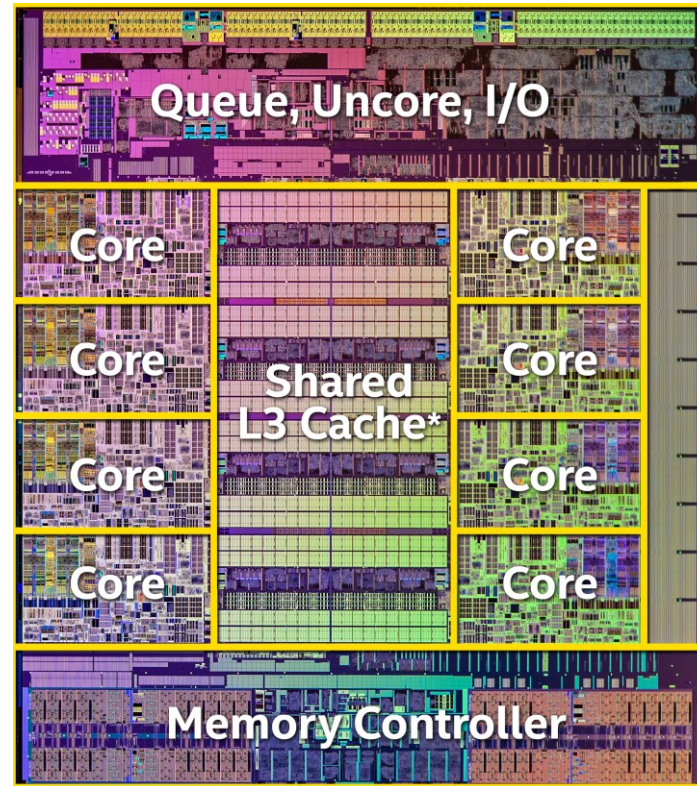
- Johannes Ahlmann
- [sensatus.io](https://sensatus.io)
  - On-Prem AI Models
  - Gathering and Enriching Web Data
  - Sales & Client Intelligence
- [webdata.org](https://webdata.org)
  - Share Libraries and Best Practices
  - Bring Data Scientists and SME Companies together
  - [ForDevelopers](#)
  - [AwesomeAvailableDatasets](#)
- Contact:  
[johannes@sensatus.io](mailto:johannes@sensatus.io)




# Times Have Changed



Intel 80286  
1 core, 0.13M transistors



Intel i7-9900K  
8 cores, 3000M transistors



# Challenges 1 / 2

Many fantastic libraries written in C, Fortran, Python, Go

How can we leverage libraries together that are written in different languages, or for different (virtual) machines?

Multiprocessing and concurrency are hard in C, Python, etc.


How can we leverage compute of multiple cores per machine, let alone multiple machines

Distributed computing is hard

How can we monitor/ manage/ debug/ reason about it?

# Challenges 2 / 2

- Data Locality
- Load Balancing
- Scaling
- State Persistence
  - Start/ Stop
  - Error Recovery



# JVM Frameworks to the Rescue

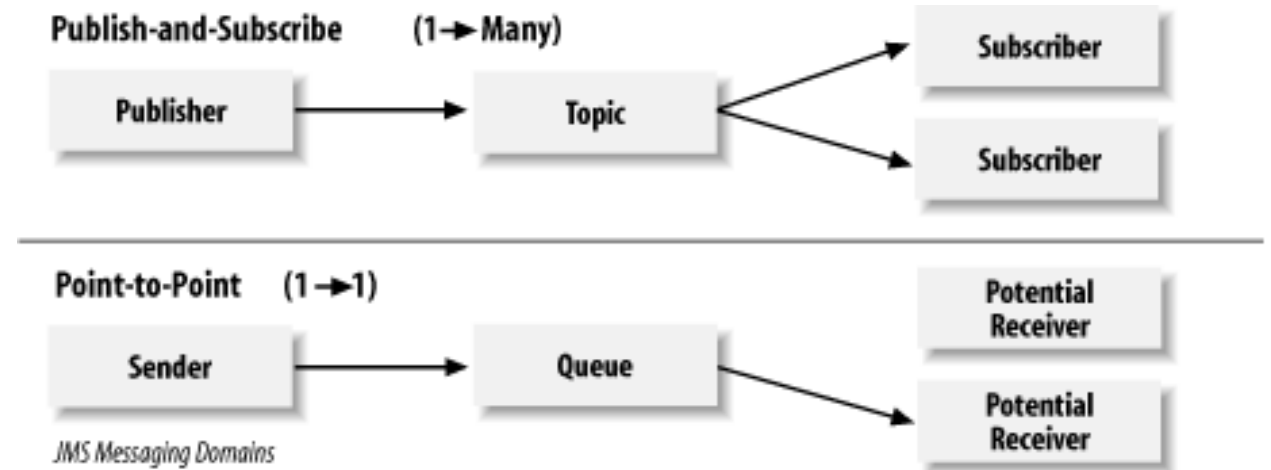






# Background: Queues & PubSub

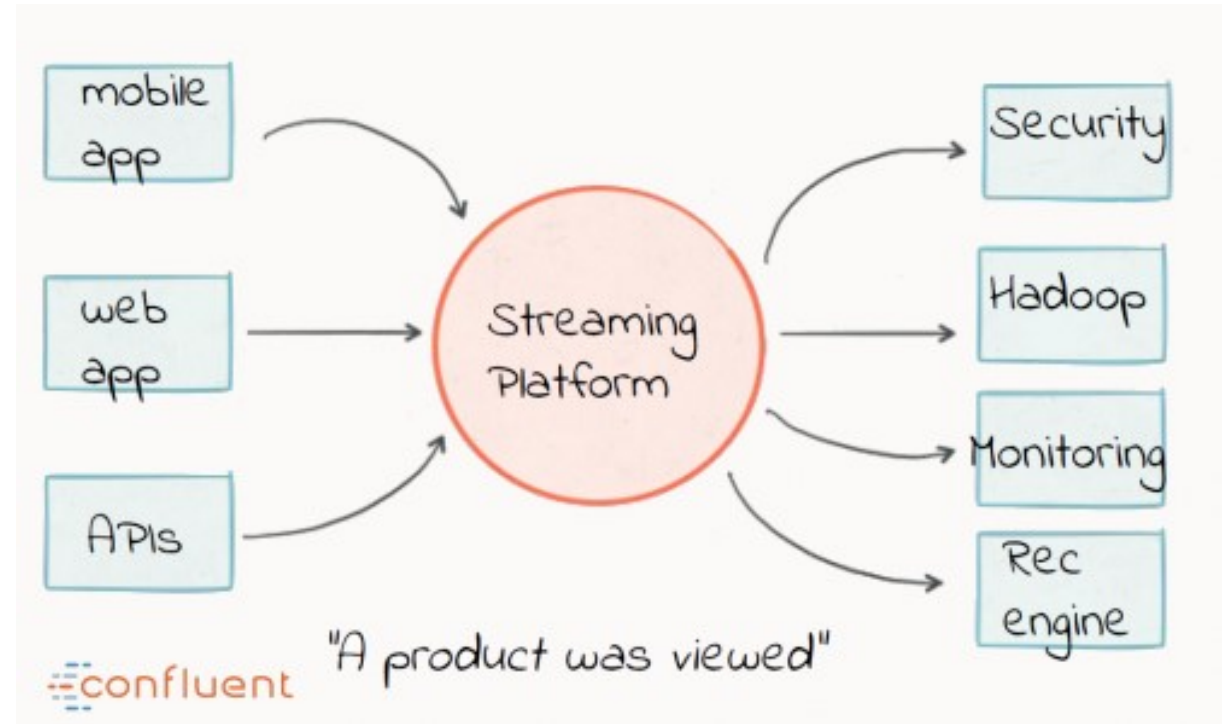
- PubSub
  - multiple subscribers
  - no scaling
- Queues
  - multiple consumers
  - single-subscriber
  - scaling, load balancing






# A high-throughput distributed messaging system

- Decouples Data Pipelines
- Scalable & Fault-Tolerant
- Kafka Functionalities
  - Messaging
  - Processing
  - Storing
- Performance (>100k/s)
  - Batching
  - Zero Copy I/O
  - Leverages OS Cache
- Durability



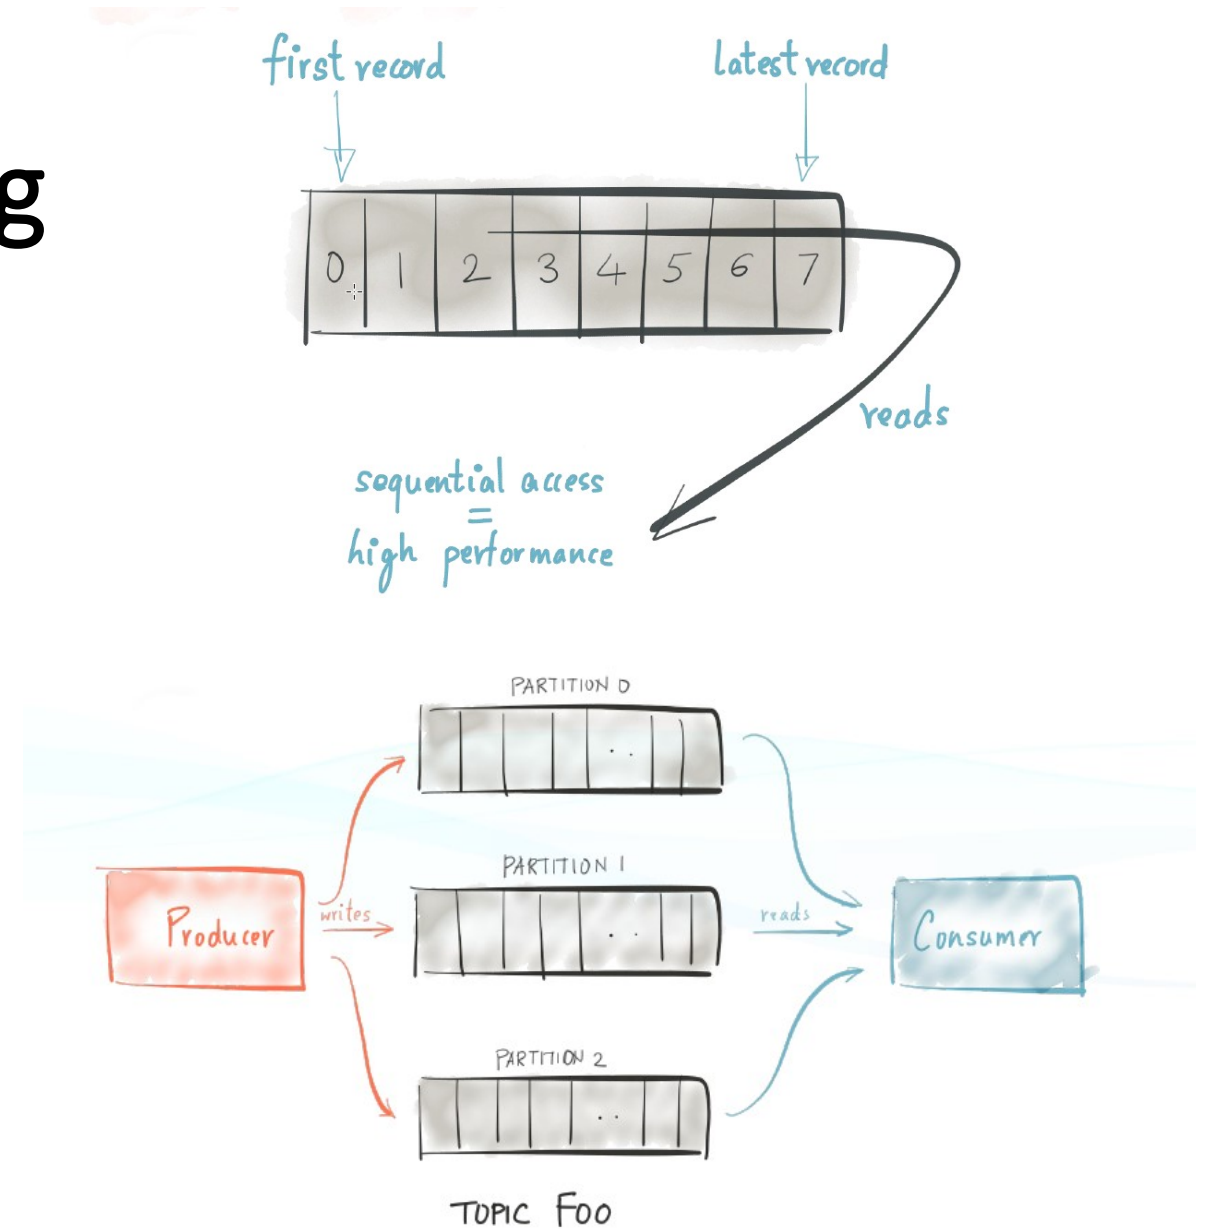
# Core Concepts

- Producers
- Consumers
- Brokers
- Topics
  - Offsets
  - Broker Addresses




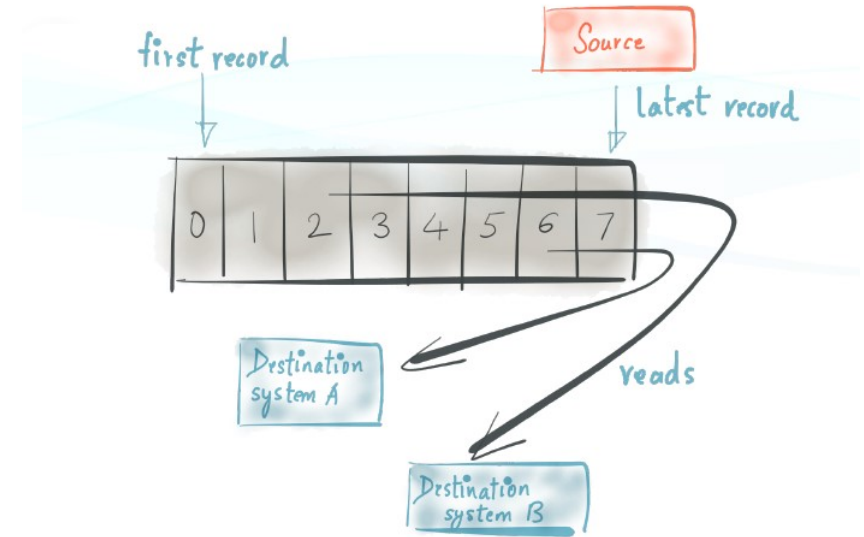
# Key Idea: Partitioned Log

- Very fast, due to zero copy I/O and batching
- Uses sendfile and OS buffer cache
- Sequential writes to FS
- Order guaranteed within partition
- Scaling



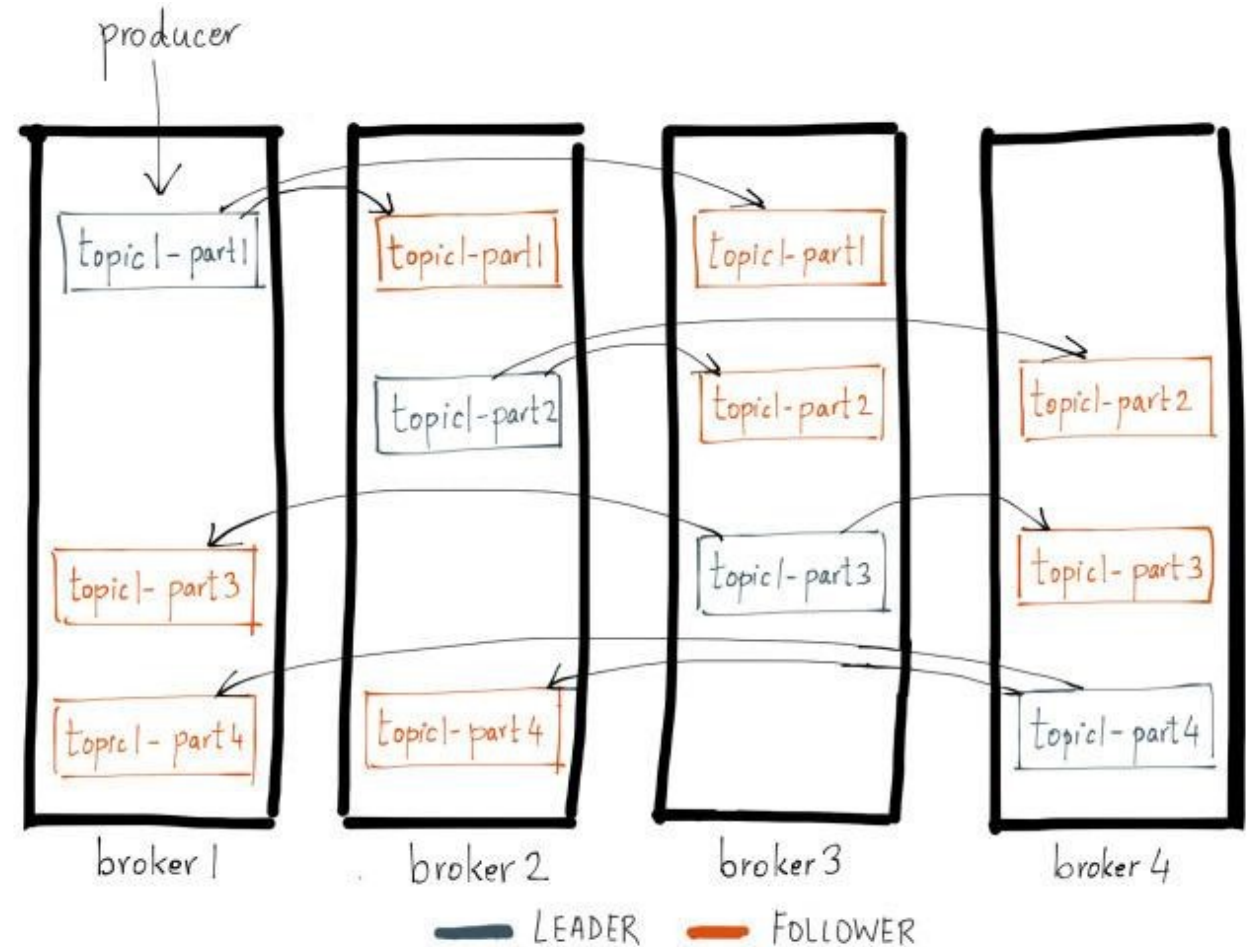
# Logs & PubSub

- Consumers can be transient
- Consumer Groups
- Delivery Semantics
  - at least once (default)
  - at most once
  - exactly once
- Retention Policy
- Reprocessing




# Partitions & Replication

- Partitions configurable
- Partition allocation
  - round-robin
  - semantic partition by key
- Replication
  - optional
  - 1 leader, 0 or more followers
  - sync or async
  - flush delay configurable



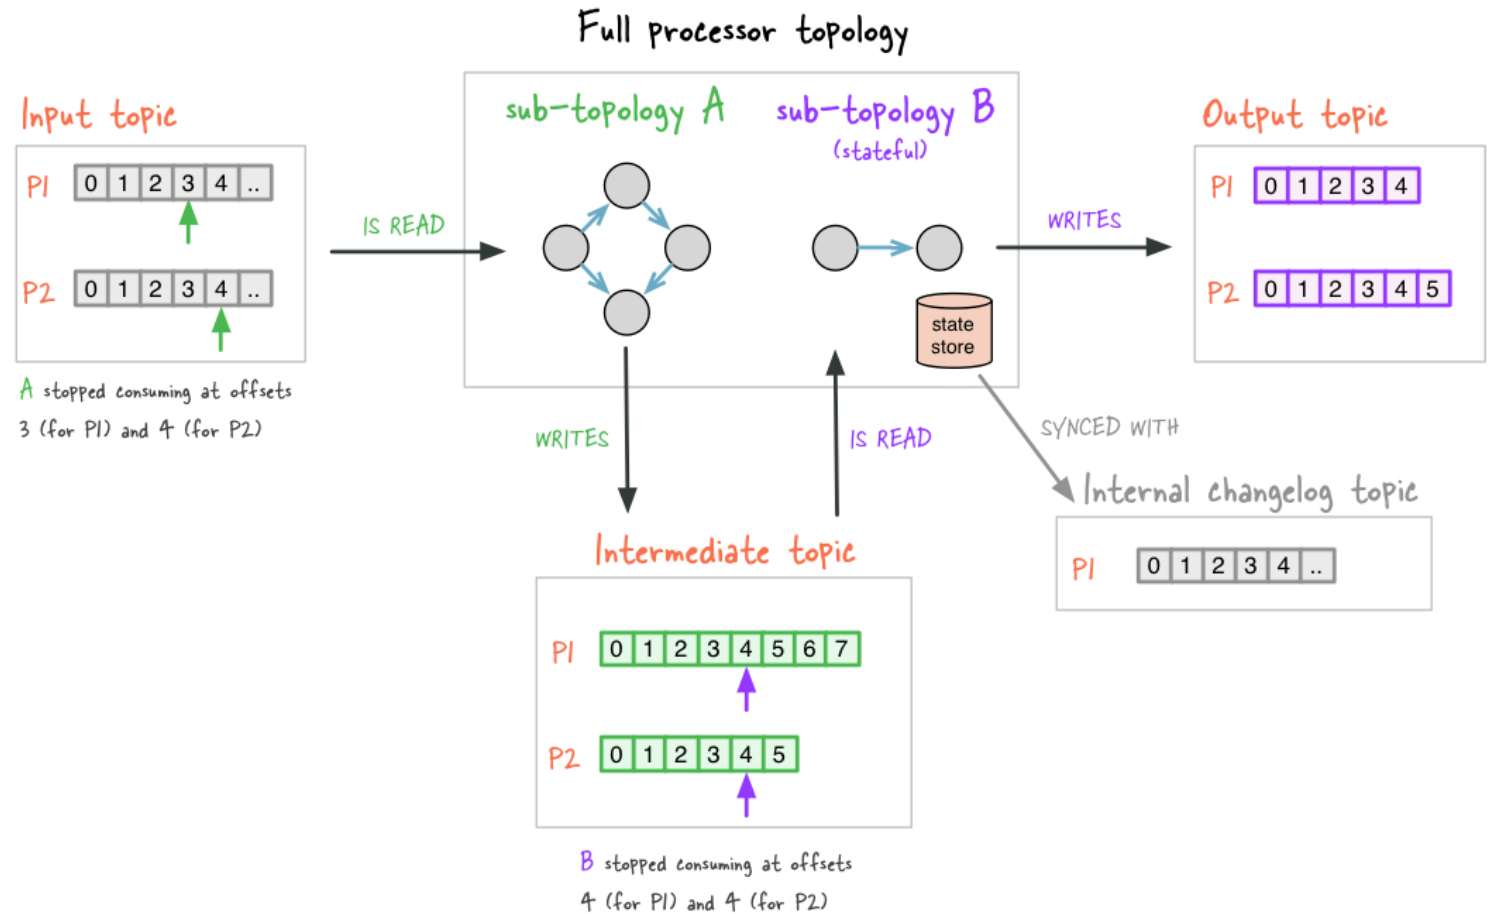
# Kafka Connect



# Kafka Streams

- Operations
  - filter
  - map
  - join
  - aggregate
- KStream
- KTable
  - manages local state
- Windows

} stateful





# Summary

## scalability of a filesystem

- hundreds of MB/s
- many TBs per server
- commodity hardware

## guarantees of a database

- persistence
- ordering


## distributed by design<sup>I</sup>

- replication
- partitioning
- horizontal scalability
- fault tolerance

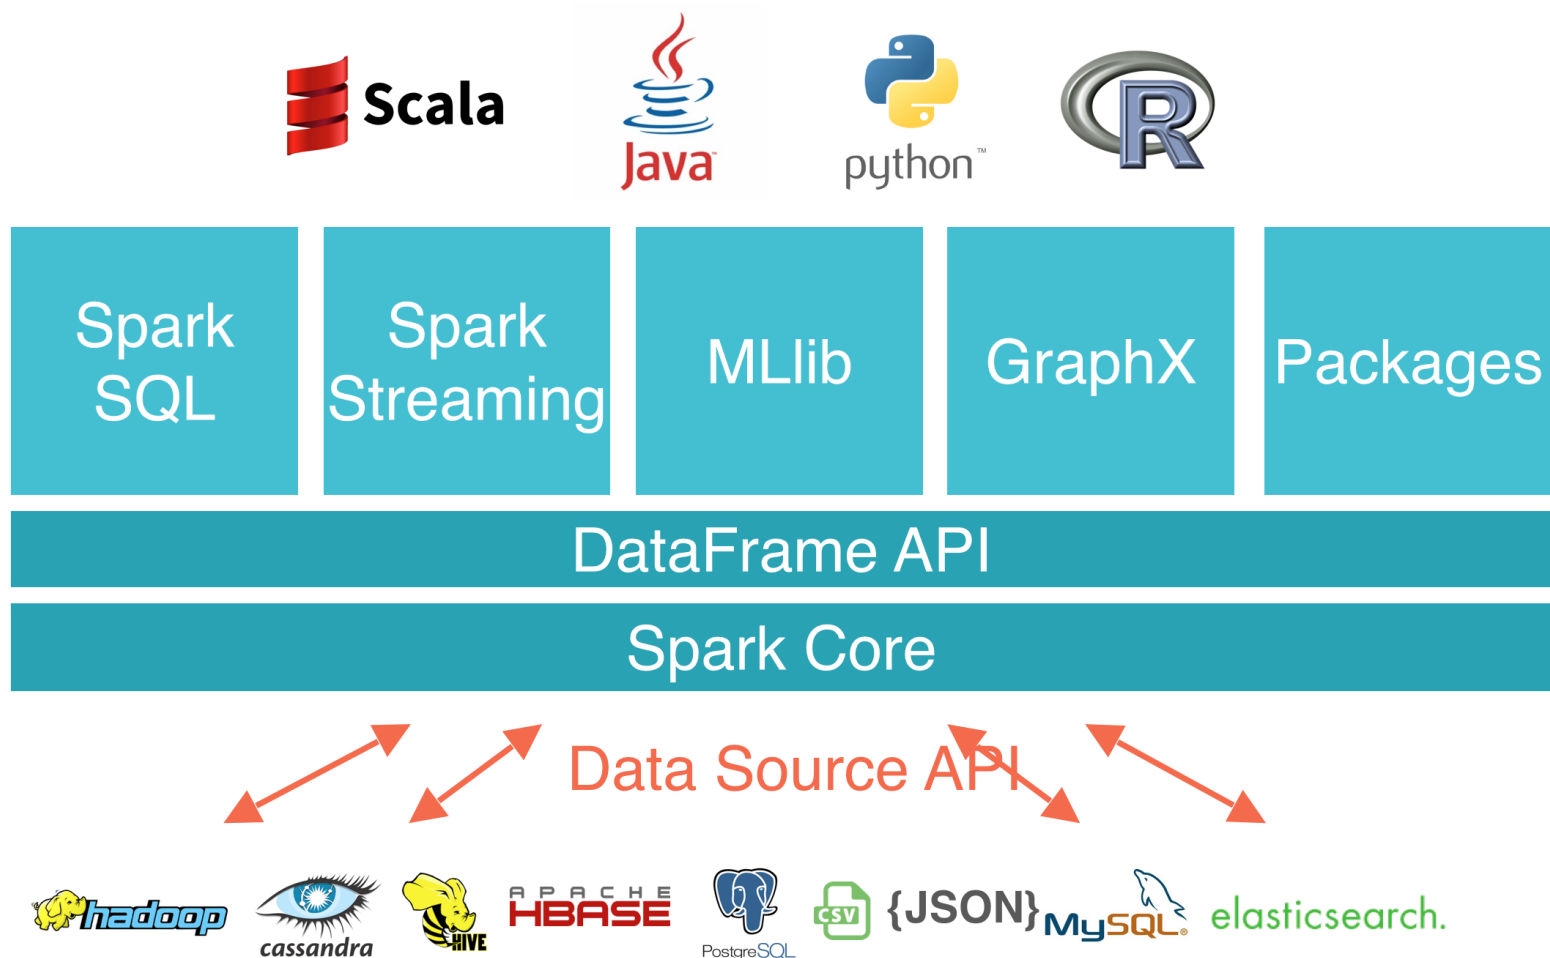


# What is Apache Spark

- Fast and general engine for large-scale data processing
- Multi-stage in-memory primitives
- Supports Iterative Algorithms
- High-Level Abstractions
- Extensible; integrated stack of libraries



# Broad Data, Language Support



# Spark Example

```
from pyspark import SparkContext
sc = SparkContext("local", "Log Analyzer")

rdd1 = sc.textFile('testfile1.txt')
rdd2 = sc.textFile('testfile2.txt')


def splitLine(line):
    key,val = line.split('|')
    return (key, int(val))

a = rdd1.filter(lambda line: not "ERROR" in line).map(splitLine)
b = rdd2.map(splitLine).reduceByKey(lambda a,b: a+b)

res = a.join(b)
#print res.toDebugString()
res.take(10)

[(u'key3', (50, 30)), (u'key1', (20, 60))]
```

# Operator Graph






vs.



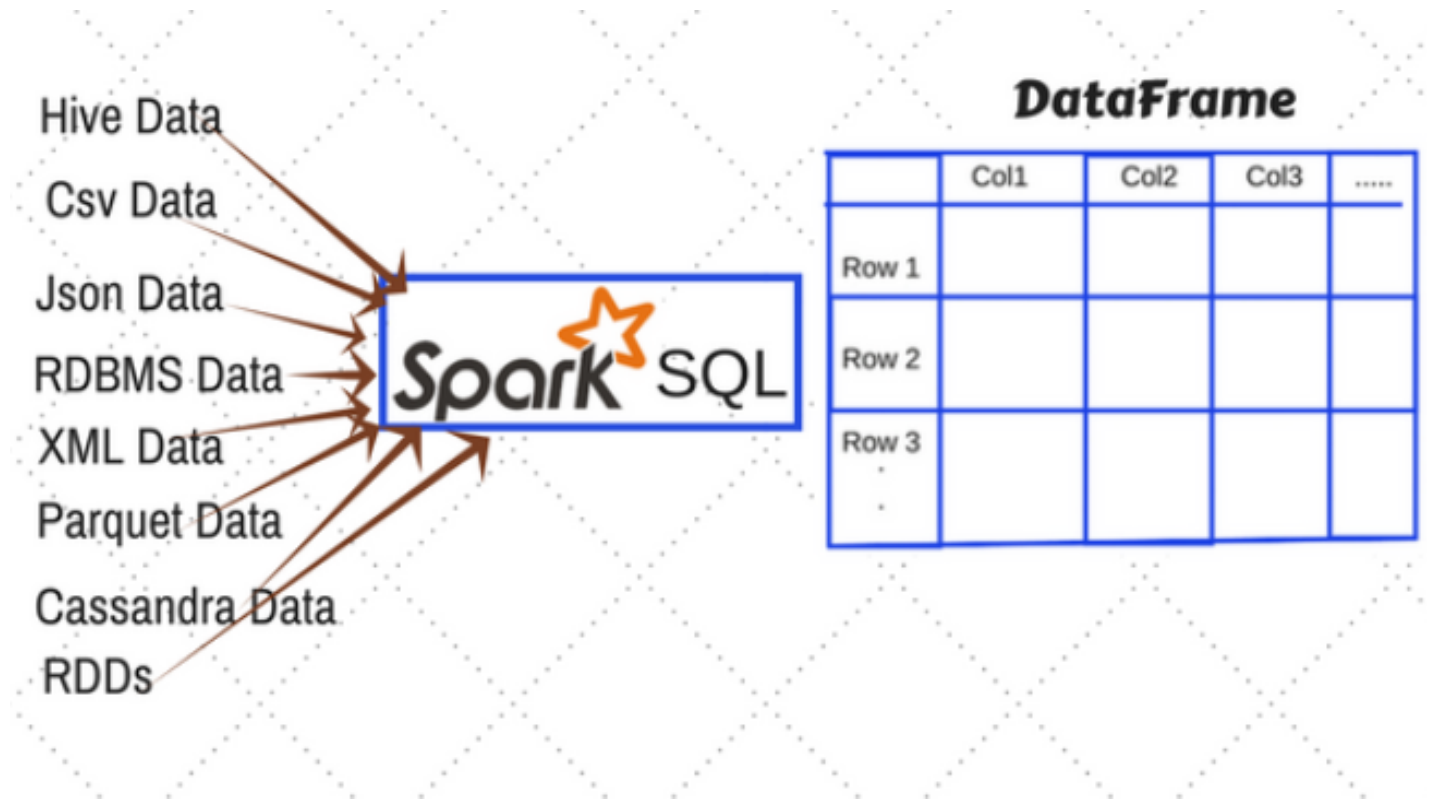
- Arbitrary operator graph
- Lazy eval of lineage graph => optimization
- Off-heap use of large memory
- Native integration with python



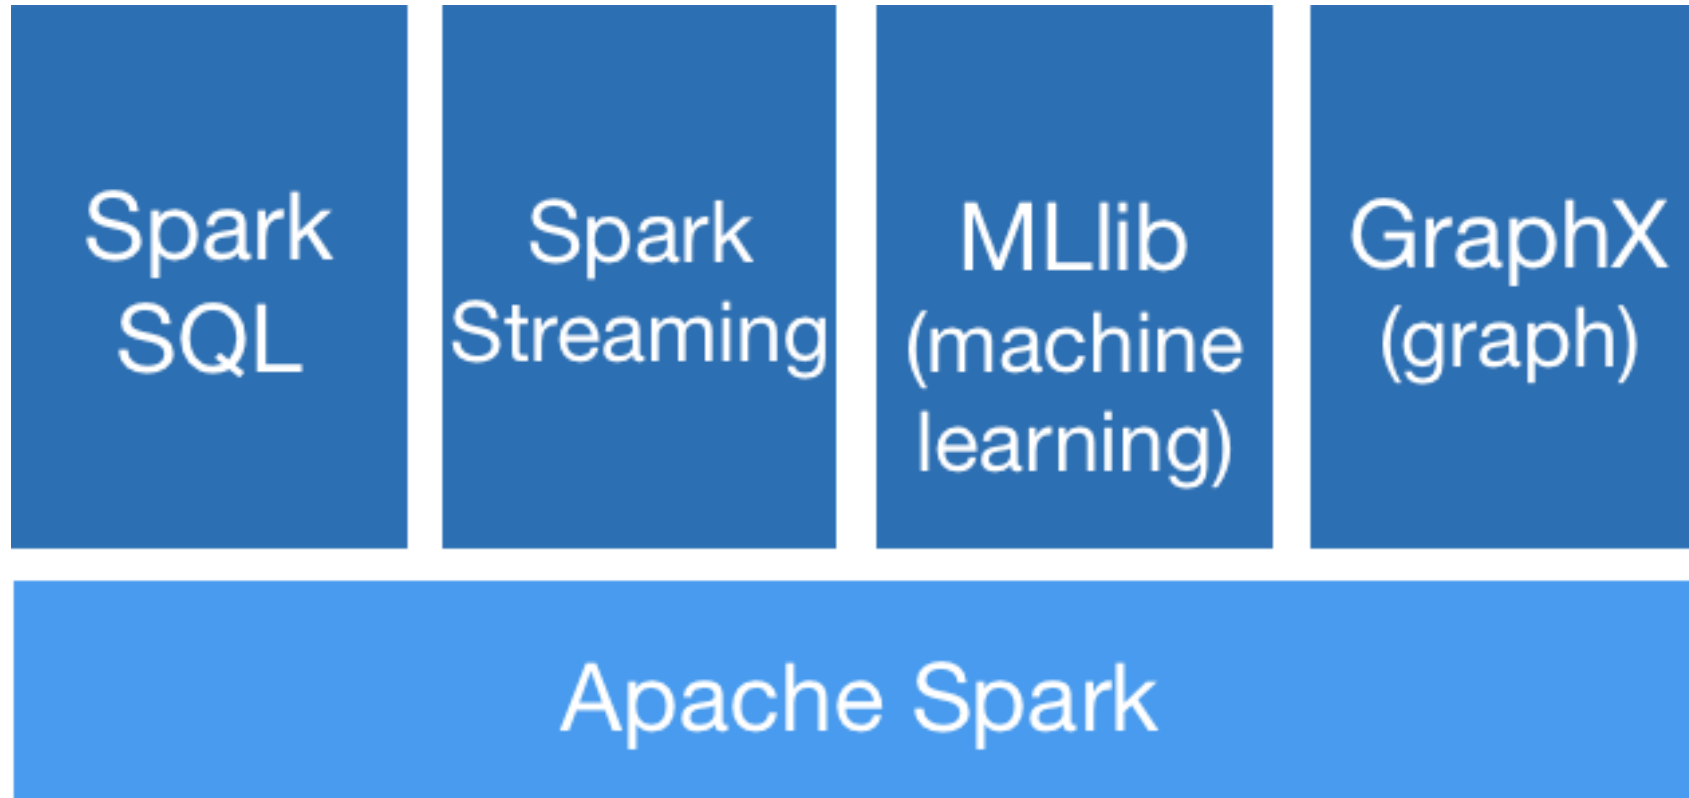


# Spark DataFrame

- Distributed Collection
- Named Columns
- Rows, Column, Schema



# Integrated Libraries



# Spark Challenges 1 / 2

**ExecutorLostFailure (executor 16 exited caused by one of the running tasks) Reason: Container killed by YARN for exceeding memory limits. 10.4 GB of 10.4 GB physical memory used. Consider boosting spark.yarn.executor.memoryOverhead.**

# Spark Challenges 2 / 2

```
17/09/05 16:40:33 ERROR Utils: Aborting task
java.lang.NullPointerException
    at org.apache.parquet.it.unimi.dsi.fastutil.objects.Object2IntLinkedOpenHashMap.getInt(Object2IntLinkedOpenHashMap.java:590)
    at org.apache.parquet.column.values.dictionary.DictionaryValuesWriter$PlainFixedLenArrayDictionaryValuesWriter.writeBytes(DictionaryValuesWriter.java:307)
    at org.apache.parquet.column.values.fallback.FallbackValuesWriter.writeBytes(FallbackValuesWriter.java:162)
    at org.apache.parquet.column.impl.ColumnWriterV1.write(ColumnWriterV1.java:201)
    at org.apache.parquet.io.MessageColumnIO$MessageColumnIORecordConsumer.addBinary(MessageColumnIO.java:467)
    at
org.apache.spark.sql.execution.datasources.parquet.ParquetWriteSupport$$anonfun$org$apache$spark$sql$execution$databases$parquet$ParquetWriteSupport$$makeWriter$10.apply(ParquetWriteSupport.scala:184)
    at
org.apache.spark.sql.execution.datasources.parquet.ParquetWriteSupport$$anonfun$org$apache$spark$sql$execution$databases$parquet$ParquetWriteSupport$$makeWriter$10.apply(ParquetWriteSupport.scala:172)
    at
org.apache.spark.sql.execution.datasources.parquet.ParquetWriteSupport$$anonfun$org$apache$spark$sql$execution$databases$parquet$ParquetWriteSupport$$writeFields$1.apply$mcV$sp(ParquetWriteSupport.scala:124)
    at
org.apache.spark.sql.execution.datasources.parquet.ParquetWriteSupport.org$apache$spark$sql$execution$databases$parquet$ParquetWriteSupport$$consumeField(ParquetWriteSupport.scala:437)
    at org.apache.spark.sql.execution.datasources.parquet.ParquetWriteSupport.org$apache$spark$sql$execution$databases$parquet$ParquetWriteSupport$$writeFields(ParquetWriteSupport.scala:123)
    at org.apache.spark.sql.execution.datasources.parquet.ParquetWriteSupport$1.apply$mcV$sp(ParquetWriteSupport.scala:114)
    at org.apache.spark.sql.execution.datasources.parquet.ParquetWriteSupport.consumeMessage(ParquetWriteSupport.scala:425)
    at org.apache.spark.sql.execution.datasources.parquet.ParquetWriteSupport.write(ParquetWriteSupport.scala:113)
    at org.apache.spark.sql.execution.datasources.parquet.ParquetWriteSupport.write(ParquetWriteSupport.scala:51)
    at org.apache.parquet.hadoop.InternalParquetRecordWriter.write(InternalParquetRecordWriter.java:123)
    at org.apache.parquet.hadoop.ParquetRecordWriter.write(ParquetRecordWriter.java:180)
    at org.apache.parquet.hadoop.ParquetRecordWriter.write(ParquetRecordWriter.java:46)
    at org.apache.spark.sql.execution.datasources.parquet.ParquetOutputWriter.write(ParquetOutputWriter.java:114)
    at org.apache.spark.sql.execution.datasources.FileFormatWriter$DynamicPartitionWriteTask.write(ParquetOutputWriter.java:114)
    at org.apache.spark.sql.execution.datasources.FileFormatWriter$DynamicPartitionWriteTask.write(ParquetOutputWriter.java:114)
    at scala.collection.Iterator$class.foreach(Iterator.scala:893)
```

***...78 more lines of  
Java stack trace***

# Takeaways

## Spark...

- feels like native python, very nice API
- adds awesome Distributed Computing and Parallel Programming capabilities to python
- comes with batteries included (SQL, GraphX, MLlib, Streaming, etc.)
- can be used from the start for exploratory programming

